

# SemLink Documentation

## Contents

<b>1</b>	<b>Prerequisites: Setting up the Conda Environment</b>	<b>2</b>
1.1	Install Miniconda . . . . .	2
1.1.1	For Linux . . . . .	2
1.1.2	For macOS . . . . .	2
1.2	Create and Activate the Environment . . . . .	2
1.3	Verify the Installation . . . . .	3
<b>2</b>	<b>Pipeline Overview</b>	<b>4</b>
2.1	Overall Pipeline Flow . . . . .	4
2.2	Step 1: Data Loading and Preprocessing ( <code>data_loader.py</code> ) . . . . .	4
2.3	Step 2: Schema Generation ( <code>schema_generator.py</code> ) . . . . .	5
2.4	Step 3: Semantic Annotation ( <code>semantic_annotation.py</code> ) . . . . .	5
2.5	Step 4: Join Discovery and Graph Generation ( <code>join_discoverer.py</code> ) . . . . .	6
<b>3</b>	<b>Datasets</b>	<b>8</b>
<b>4</b>	<b>Related Repositories</b>	<b>9</b>

# 1 Prerequisites: Setting up the Conda Environment

This pipeline requires a dedicated Python environment with a specific set of libraries. The easiest way to manage this is with Conda. If you don't have Conda installed, you can follow the steps below.

## 1.1 Install Miniconda

Miniconda is a minimal installer for Conda, Python, and their core dependencies. It is the recommended way to get started.

### 1.1.1 For Linux

Use `wget` or `curl` to download the installer script, then run it.

```
# Download the latest Linux installer
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
-O miniconda.sh

# Run the installer in silent mode (-b) and force a specific prefix (-p)
bash miniconda.sh -b -p $HOME/miniconda3

# Remove the installer script
rm miniconda.sh

# Initialize Conda for your shell (e.g., bash or zsh)
~/miniconda3/bin/conda init
```

### 1.1.2 For macOS

Use `curl` to download the installer script, then run it. For Apple Silicon Macs, use the `arm64` installer. For Intel-based Macs, use the `x86_64` installer.

```
# For Apple Silicon (M1/M2/M3)
curl https://repo.anaconda.com/miniconda/Miniconda3-latest-MacOSX-arm64.sh
-o miniconda.sh

# For Intel-based Macs
curl https://repo.anaconda.com/miniconda/Miniconda3-latest-MacOSX-x86_64.sh
-o miniconda.sh

# Run the installer in silent mode (-b) and force a specific prefix (-p)
bash miniconda.sh -b -p $HOME/miniconda3

# Remove the installer script
rm miniconda.sh

# Initialize Conda for your shell (e.g., zsh)
~/miniconda3/bin/conda init zsh
```

After running `conda init`, you may need to close and reopen your terminal for the changes to take effect. You should see `(base)` at the beginning of your prompt, indicating the base Conda environment is active.

## 1.2 Create and Activate the Environment

Once Miniconda is installed, navigate to the directory containing your `environment.yml` file and run the following command to create the environment. This will install all the required packages listed in the file.

```
conda env create -f environment.yml
```

Once the environment is created, activate it with the following command:

```
conda activate semlink
```

### **1.3 Verify the Installation**

To ensure everything is installed correctly, you can list the packages in your new environment:

```
conda list
```

## 2 Pipeline Overview

This pipeline offers a structured way to load raw data, semantically annotate it, and discover potential join relationships. It uses a series of modules to transform raw data, enrich it with meaningful metadata, and create a graph-based model of its relationships.

### 2.1 Overall Pipeline Flow

The pipeline consists of four main steps, executed in order to achieve the final output:

1. **Data Loading and Preprocessing:** The `data_loader.py` module takes a directory of CSV/TSV files and converts them into a standardized JSON format. It performs basic statistical analysis and, if a Large Language Model (LLM) is used, generates initial descriptions.
2. **Schema Generation:** The `schema_generator.py` module uses the processed data to generate a LinkML schema. This schema acts as a conceptual model, defining classes and attributes that represent the data. This step also leverages an LLM to suggest a schema based on column names and descriptions.
3. **Semantic Annotation:** The `semantic_annotation.py` module takes the processed data and the generated LinkML schema to semantically annotate each column. It maps each column to a specific class and attribute within the schema.
4. **Join Discovery and Graph Generation:** The final step, handled by `join_discoverer.py`, uses the semantic annotations to find potential joinable columns. It generates embeddings for each column and computes similarity metrics. The final output is a set of CSV files for use in a graph database like Neo4j, representing a graph of data nodes and their relationships.

### 2.2 Step 1: Data Loading and Preprocessing (`data_loader.py`)

This is the starting point of the pipeline. It handles ingesting your raw data and generating an initial structured representation.

- **Input:**
  - A directory containing `.csv` or `.tsv` files.
  - (Optional) A directory containing `.json` metadata files corresponding to the `.csv` files.
  - (Optional) An initialized OpenAI client for LLM-based descriptions.
- **Input Data Structure:**
  - The pipeline expects semi-structured data in CSV or TSV format.
  - Each file is treated as a single table, with column headers used to identify columns.
- **Output:**
  - A JSON file named `data_lake.json`.
  - The output is a list of dictionaries, where each dictionary represents a file (table) and contains metadata about the file and its columns. This includes basic statistics and, if an LLM is used, high-level descriptions.

Here's how to run this step:

```
# Import the necessary class
from data_loader import DataLoader
from openai import OpenAI

# Initialize OpenAI client and DataLoader
client = OpenAI(api_key="<YOUR_OPENAI_API_KEY>")
loader = DataLoader(openai_client=client)

# Run the data loading and processing step
```

```
data_lake = loader.load_and_describe_datalake(
    directory='<YOUR.DATA.DIRECTORY.PATH>',
    output_directory='<YOUR.OUTPUT.DIRECTORY.PATH>'
)
```

## 2.3 Step 2: Schema Generation (schema\_generator.py)

This step creates a conceptual schema from your data, which is essential for semantic annotation.

- **Input:**

- The `data_lake.json` file generated in Step 1.
- An initialized OpenAI client.
- (Optional) A pre-existing LinkML schema file (`.yaml`) for pruning only.

- **Input Data Structure:**

- The primary input is the `data_lake.json`. The module uses the `llm_description` and `llm_column_description` fields to inform the schema generation.

- **Output:**

- A LinkML schema file named `linkml_schema.yaml`.
- A pruned version of the schema named `linkml_schema_pruned.yaml`.
- This schema defines classes (entities) and attributes (their properties), providing a standardized vocabulary.

Here's how to run this step:

```
# Import the necessary class
from schema_generator import SchemaGenerator

# Initialize SchemaGenerator
schema_gen = SchemaGenerator(openai_client=client)

# Run the schema generation step
linkml_schema = schema_gen.generate_schema_from_datalake(
    data_lake=data_lake,
    output_directory='<YOUR.OUTPUT.DIRECTORY.PATH>'
)

# Prune the schema
pruned_schema = schema_gen.prune_schema(
    yaml_schema=linkml_schema,
    data_lake=data_lake,
    output_directory='<YOUR.OUTPUT.DIRECTORY.PATH>'
)
```

## 2.4 Step 3: Semantic Annotation (semantic\_annotation.py)

This step enriches your data by mapping it to the generated schema, adding a semantic layer that describes the meaning of each column.

- **Input:**

- The `data_lake.json` file from Step 1.
- The `linkml_schema_pruned.yaml` file from Step 2.

- **Input Data Structure:**

- The module takes the `data_lake.json` and uses the column names and sample values from each file.
- The `linkml_schema_pruned.yaml` file provides the target vocabulary for annotation.

- **Output:**

- An annotated JSON file named `data_lake_annotated.json`.
- The output is the same as `data_lake.json`, but with a new `semantic_annotation` field for each column, specifying its corresponding `class.attribute` from the LinkML schema.

Here's how to run this step:

```
# Import the necessary function
from semantic_annotation import archetype_annotation

# Run the semantic annotation step
data_lake_annotated = archetype_annotation(
    data_lake_list=data_lake,
    yaml_schema=pruned_schema,
    output_directory='<YOUR_OUTPUT_DIRECTORY_PATH>'
)
```

## 2.5 Step 4: Join Discovery and Graph Generation (`join_discoverer.py`)

The final step analyzes the semantically annotated data to identify and visualize potential join points.

- **Input:**

- The `data_lake_annotated.json` file from Step 3.
- (Optional) A pre-computed embeddings JSON file to skip the embedding generation.
- An initialized OpenAI client.

- **Input Data Structure:**

- The module requires the `semantic_annotation` field in the input JSON to be populated. It uses this information to generate unique identifiers for each column.
- The OpenAI client generates vector embeddings for each column's context, which are used for similarity calculations.

- **Output:**

- A JSON file named `embeddings.json` containing the vector embeddings for each column.
- A set of CSV files (`distances.csv`, `nodes.csv`, and `edges.csv`) for importing into a graph database.
- `nodes.csv`: Represents each column as a node.
- `edges.csv`: Represents discovered join relationships as edges between nodes, with a `cosine_similarity` property.
- `distances.csv`: A detailed breakdown of all computed similarity and distance values.

Here's how to run this step:

```
# Import the necessary class
from join_discoverer import JoinDiscoverer

# Initialize JoinDiscoverer
join_dis = JoinDiscoverer(openai_client=client)

# Generate embeddings
embeddings = join_dis.generate_embeddings(
    data_lake_annotated=data_lake_annotated,
```

```

        output_directory='<YOUR_OUTPUT_DIRECTORY_PATH>'
    )
# Compute distances and export to Neo4j-compatible CSVs
-, -, - = join_dis.compute_distances_and_export_neo4j(
    embeddings=embeddings,
    cosine_sim_threshold=0.8,
    anns_threshold=1.5,
    output_directory='<YOUR_OUTPUT_DIRECTORY_PATH>'
)

```

### 3 Datasets

- Valentine
- PheKnowLator benchmark KG
- Eurostat



## 4 Related Repositories

**ArcheType** – Module for semantic annotation. To download and place in the Src folder.